


Operating-System Debugging

- **Debugging** is finding and fixing errors, or **bugs**
- OS generate **log files** containing error information
- Failure of an application can generate **crash dump** file capturing memory of the process.
 - Location may be found or modified by editing the file
`/proc/sys/kernel/core_pattern`
- Operating system failure can generate **crash dump** file containing kernel memory
 - `/var/crash/vmcore`, but configured in `/etc/kdump.conf`
- Beyond crashes, performance tuning can optimize system performance
 - Sometimes using ***trace listings*** of activities, recorded for analysis
 - **Profiling** is periodic sampling of instruction pointer (i.e. program counter, PC) to look for statistical trends.
 - **perf** is a userspace utility that is accessed from the command line and provides a number of subcommands; it is capable of statistical profiling of the entire system (both kernel and user code).

Performance Tuning



- Improve performance by removing bottlenecks
- OS must provide means of computing and displaying measures of system behavior
- Windows: **Task manager**
- Linux (Gnome or KDE): “**system monitor**”

Performance Tuning

Activities System Monitor ▾

Wed 14:16

Processes Resources File Systems


The screenshot shows the System Monitor application window. The title bar includes the application name, a search icon, and system status icons. The main area displays a table of processes with columns for Process Name, User, % CPU, ID, Memory, Disk read total, Disk write t, Disk read, Disk write, Nice, and Priority. A horizontal bar at the top of the table indicates the current selection. The table lists various processes, including system daemons like kswapd0, systemd, and ksoftirqd, along with user applications like gnome-system-monitor, gnome-shell, Xorg, whoopsie, and rsyslogd. The row for kswapd0 is highlighted with an orange background.

Process Name	User	% CPU	ID	Memory	Disk read total	Disk write t	Disk read	Disk write	Nice	Priority
gnome-system-monitor	omansour		1	6802 13.2 MiB	49.8 MiB	48.0 KiB	N/A	N/A	N/A	0 Normal
gnome-shell	omansour		0	1924 218.1 MiB	N/A	N/A	N/A	N/A	N/A	0 Normal
Xorg	omansour		0	1791 43.7 MiB	N/A	N/A	N/A	N/A	N/A	0 Normal
whoopsie	whoopsie		0	1090 992.0 KiB	N/A	N/A	N/A	N/A	N/A	0 Normal
systemd-resolved	systemd-re:		0	458 236.0 KiB	N/A	N/A	N/A	N/A	N/A	0 Normal
rsyslogd	syslog		0	714 428.0 KiB	N/A	N/A	N/A	N/A	N/A	0 Normal
rtkit-daemon	rtkit		0	1209 N/A	N/A	N/A	N/A	N/A	N/A	1 Normal
kswapd0	root		0	55 N/A	N/A	N/A	N/A	N/A	N/A	0 Normal
systemd	root		0	1 2.4 MiB	N/A	N/A	N/A	N/A	N/A	0 Normal
kthreadd	root		0	2 N/A	N/A	N/A	N/A	N/A	N/A	0 Normal
kworker/0:0H	root		0	4 N/A	N/A	N/A	N/A	N/A	N/A	-20 Very High
mm_percpu_wq	root		0	6 N/A	N/A	N/A	N/A	N/A	N/A	-20 Very High
ksoftirqd/0	root		0	7 N/A	N/A	N/A	N/A	N/A	N/A	0 Normal
rcu_sched	root		0	8 N/A	N/A	N/A	N/A	N/A	N/A	0 Normal
rcu_bh	root		0	9 N/A	N/A	N/A	N/A	N/A	N/A	0 Normal
migration/0	root		0	10 N/A	N/A	N/A	N/A	N/A	N/A	0 Normal
watchdog/0	root		0	11 N/A	N/A	N/A	N/A	N/A	N/A	0 Normal
cpuhp/0	root		0	12 N/A	N/A	N/A	N/A	N/A	N/A	0 Normal
cpuhp/1	root		0	13 N/A	N/A	N/A	N/A	N/A	N/A	0 Normal
watchdog/1	root		0	14 N/A	N/A	N/A	N/A	N/A	N/A	0 Normal
migration/1	root		0	15 N/A	N/A	N/A	N/A	N/A	N/A	0 Normal
ksoftirqd/1	root		0	16 N/A	N/A	N/A	N/A	N/A	N/A	0 Normal
kworker/1:0H	root		0	18 N/A	N/A	N/A	N/A	N/A	N/A	-20 Very High
cpuhp/2	root		0	19 N/A	N/A	N/A	N/A	N/A	N/A	0 Normal
watchdog/2	root		0	20 N/A	N/A	N/A	N/A	N/A	N/A	0 Normal
migration/2	root		0	21 N/A	N/A	N/A	N/A	N/A	N/A	0 Normal
ksoftirqd/2	root		0	22 N/A	N/A	N/A	N/A	N/A	N/A	0 Normal
kworker/2:0H	root		0	24 N/A	N/A	N/A	N/A	N/A	N/A	-20 Very High

End Process

Virtualization

- Allows operating systems to run applications within other Oses
 - Vast and growing industry
- A **host** system with an OS and a VMM may run one or more **guest** systems, each with a different OS
- Some VMM's may install directly on Hardware, without a need for a host OS.



Virtualization – cont.

- **Emulation** used when guest system is compiled for a CPU type that is different from that of the host.
 - **Interpretation** is used to map guest machine instructions to host CPU's machine instructions. E.g. Apple's Parallels
 - Generally slow.
 - e.g. old Apple software (PowerPC CPUs) running on newer Apple machines (x86 CPUs).
 - e.g. Running a windows system on Apple machines prior to 2006 (PowerPC CPU's) or on Apple machines running M1/M2 CPU's.
- **Virtualization** – Host OS natively compiled for CPU architecture (always the case), running guest OSes also natively compiled for same CPU instruction set architecture.
 - **VMM** (virtual machine Manager) provides virtualization services
 - Consider a Windows 10 **host** OS running running Windows 10 guests (via Vmware VMM)
 - Or, a Windows 10 host running Ubuntu Linux (via Vmware or VirtualBox VMM)

Virtualization - cont.

- **Use cases** involve laptops, desktops or servers running multiple OSes for **exploration or compatibility**
 - Apple laptop running Mac OS X host, Windows as a guest: e.g. parallels “Desktop” and VMware’s “Fusion” allows a Mac system to run windows and windows applications, or Linux.
 - Windows laptop with Linux as a guest: We shall use this setup for our labs, using VMware workstation.
 - **Developing apps** for multiple OSes without having multiple systems
 - **QA testing** applications without having multiple systems
 - Executing and managing compute environments **within data centers**
- A VMM may also run without a host OS (Type 2 Hypervisor)
 - VMM replaces the general-purpose OS and becomes the host OS.
 - e.g. VMware ESX and Citrix XenServer

Open-Source Operating Systems


- Operating systems made available in source-code format rather than just binary **closed-source**
- Examples include:
 - **GNU-Linux** and
 - **Android-Linux**
 - **BSD UNIX** (e.g. FreeBSD and NetBSD)
 - and a few more.
- When building or developing applications and drivers for an open source OS, good attention must be paid to the licensing model:
 - Reciprocal: e.g. GPLv1, GPLv2, etc.
 - Permissive: e.g. BSD, MIT, etc.

Review

- C – review
 - Function pointers
- Linking and kernel modules
- Elementary data structures - review


C-language review - Pointers

- Declaring a C variable allocates a number of memory cells (or bytes) and assigns them a name (the name of the variable).
- Memory is organized into bytes (or cells), where each has a unique address.
- When a variable is allocated a certain number of bytes in memory, they are always contiguous, for example:
`int x;`



Pointers (cont.)

- In many cases, your program may need to know the memory address of your variable and may also need to access a variable using its address instead of its name.
 - Note: An address is always a byte-address (e.g. address 6000 tells you integer x is preceded by 6000 bytes before it, NOT 6000 integers, 4-bytes each).




Pointers – The address of operator &

- You can access the address of a variable using the & operator, e.g.

```
char x = 37;  
char *px;  
px = &x;
```

In the above example:

- & is the “**address-of**” operator
- **char* is the declaration** of a pointer




Pointers – The dereference operator *

The **dereference operator *** allows reads or writes to a variable using its pointer instead of its name, for example:


```
*px = 93;
```

change the contents of variable x from 37 to 93. Hence it is equivalent to the statement:

```
x=93;
```



Memory
before the
statement
`*px=93;`




Memory
after the
statement
`*px=93;`

Pointers – cont.

```
char x = 37;  
char *px = &x;
```

Which one of the following statements evaluates to true?

- (x==5000)
- (x==37)
- (&x==9500)
- (&x==5000)
- (px==9500)
- (px==37)
- (*px==9500)
- (*px==37)




Pointers - Example:

```
char x = 37;  
char *px = &x;
```

Which one of the following statements evaluates to true?

(x==5000)	false
(x==37)	true
(&x==9500)	false
(&x==5000)	true
(px==9500)	false
(px==37)	false
(*px==9500)	false
(*px==37)	true



Pointers - declaration

```
int *p1;  
char *p2;  
double *p3;
```

- Note that the asterisk (*) used when declaring a pointer should not be confused with the dereference operator seen earlier. They are two different things represented with the same sign.

```
int * x, y;
```

- In the previous line, x is declared as a pointer, but y is declared as an int.

Pointers – example 1

```
// my first pointer
#include <stdio.h>

int main ()
{
    int firstvalue, secondvalue;
    int * mypointer;

    mypointer = &firstvalue;
    *mypointer = 10;
    mypointer = &secondvalue;
    *mypointer = 20;

    printf("firstvalue is %d\n", firstvalue);
    printf("secondvalue is %d\n", secondvalue);
    return 0;
}
```

Pointers – example 1

```
// my first pointer
#include <stdio.h>

int main ()
{
    int firstvalue, secondvalue;
    int * mypointer;

    mypointer = &firstvalue;
    *mypointer = 10;
    mypointer = &secondvalue;
    *mypointer = 20;

    printf("firstvalue is %d\n", firstvalue);
    printf("secondvalue is %d\n", secondvalue);
    return 0;
}
```

```
firstvalue is 10
secondvalue is 20
```

Pointers – example 2

```
// more pointers
#include <stdio.h>

int main ()
{
    int firstvalue = 5, secondvalue = 15;
    int * p1, * p2;

    p1 = &firstvalue; // p1 = address of firstvalue
    p2 = &secondvalue; // p2 = address of secondvalue
    *p1 = 10;           // value pointed to by p1 = 10
    *p2 = *p1;          // value pointed to by p2 =
                        //value pointed to by p1
    p1 = p2;            // (value of pointer is copied)
    *p1 = 20;            // value pointed to by p1 = 20

    printf("firstvalue is %d\n", firstvalue);
    printf("secondvalue is %d\n", secondvalue);
    return 0;
}
```

Pointers – example 2

```
// more pointers
#include <stdio.h>

int main ()
{
    int firstvalue = 5, secondvalue = 15;
    int * p1, * p2;

    p1 = &firstvalue; // p1 = address of firstvalue
    p2 = &secondvalue; // p2 = address of secondvalue
    *p1 = 10;           // value pointed to by p1 = 10
    *p2 = *p1;          // value pointed to by p2 =
                        //value pointed to by p1
    p1 = p2;            // (value of pointer is copied)
    *p1 = 20;            // value pointed to by p1 = 20

    printf("firstvalue is %d\n",firstvalue);
    printf("secondvalue is %d\n",secondvalue);
    return 0;
}
```

```
firstvalue is 10
secondvalue is 20
```

Arrays and pointers

- The array name holds the starting address of the array

```
int vals[] = { 4, 7, 11 };
```



starting address of `vals`: 0x4a00

```
printf("%lx", (unsigned long) vals);  
        // displays 0x4a00  
printf("%lx", (unsigned long) vals[0]);  
        // displays 0x4
```

Arrays and pointers – cont.

- Array name can be used as a pointer (a **constant pointer**):

```
int vals[] = {4, 7, 11};  
printf("%d", *vals); // displays 4
```

- Pointer can be used as an array name:

```
int *valptr = vals;  
printf("%d", valptr[1]); // displays 7  
printf("%d", *valptr); // displays ??  
printf("%d", valptr[0]); // displays ??  
printf("%d", *(valptr+1)); // displays ??
```

Arrays and pointers

- Hence, arrays work very much like pointers to their first element, and an array can always be implicitly converted to a pointer of the proper type, i.e. a **pointer can be assigned any value, whereas an array can only represent the same elements it pointed to during its instantiation**, hence:

```
int x[20];  
int *px;      valid  
  
px = x;
```

~~x = px;~~ Not valid

- An array declaration allocates memory for the number of elements inside the array, whereas the declaration of a pointer allocates only the memory required to hold an address.

Arrays and pointers – example

```
// more pointers
#include <stdio.h>

int main ()
{
    int numbers[5];
    int * p;
    p = numbers;  *p = 10;
    p++;  *p = 20;
    p = &numbers[2];  *p = 30;
    p = numbers + 3;  *p = 40;
    p = numbers;  *(p+4) = 50;
    for (int n=0; n<5; n++)
        printf("%d, ", numbers[n]);
    return 0;
}
```

Array and pointers – example

```
// more pointers  
#include <stdio.h>
```

```
int main ()  
{  
    int numbers[5];  
    int * p;  
    p = numbers;  *p = 10;  
    p++;  *p = 20;  
    p = &numbers[2];  *p = 30;  
    p = numbers + 3;  *p = 40;  
    p = numbers;  *(p+4) = 50;  
    for (int n=0; n<5; n++)  
        printf("%d, ", numbers[n]);  
    return 0;  
}
```

```
10, 20, 30, 40, 50,
```

Pointers to functions

- C doesn't require that pointers point only to data;
- Function pointers point to memory addresses where functions are stored, e.g.

```
void (*fp) (void);
```

- A function's name may be viewed as a constant pointer to a function.

Pointers to functions

```
#include <stdio.h>

void func_0() {
    printf("I am func_0():\n");
}

int main() {
    void (*fp) () = func_0;
    fp();
    (*fp)();
    (fp)();
    return 0;
}
```

```
I am func_0():
I am func_0():
I am func_0():
```

- Function name may be viewed as a const pointer.
- Parameters and return type must match
- Either form may be used to invoke the function

Arrays of Pointers to functions

```
#include <stdio.h>

void func_0(){ printf("I am func_0():\n"); }
void func_1(){ printf("I am func_1():\n"); }
void func_2(){ printf("I am func_2():\n"); }
void func_3(){ printf("I am func_3():\n"); }

int main() {
    void (*fp[5])() = {
        func_0,
        func_1,
        func_2,
        func_3
    };
    fp[4] = 0;

    int i=0;
    while (fp[i]) {
        fp[i]();
        i++;
    }
    return 0;
}
```

I am func_0()
I am func_1()
I am func_2()
I am func_3()

- We can also have arrays of function pointers
- In this particular example, we are using a “0” as a sentinel.

Pointers to functions

```
#include <stdio.h>

void func_0() {
    printf("I am func_0():\n");
}

struct Cobj{
    char *name;
    void (*printname)();
};

int main() {
    // create a c-style object
    struct Cobj myObj;

    // initialize the c-string and the methods
    myObj.name = "Sam Smith";
    myObj.printname = func_0;

    // Invoke the method
    myObj.printname();


    return 0;
}
```

I am func_0():

- Function pointers may be also declared inside structs
- They may be invoked like any method you would invoke in C++ or in Java
- **This how Linux implements object-oriented concepts, despite using C, a non object-oriented language!**

The Queue data structure: (first in, first out – FIFO)

- Queue: a FIFO (first in, first out) data structure.
- Examples:
 - print jobs sent to a printer
 - Input from a keyboard is buffered into a stream using a fixed size FIFO.
 - TCP/IP packets waiting to be transmitted
- Implementation:
 - static: fixed size, implemented as array
 - dynamic: variable size, implemented as linked list



The Queue data structure - operations


- Locations
 - Back/Rear (tail): position where elements are added
 - Front (head): position from which elements are removed
- Operations:
 - enqueue: add an element to the rear of the queue
 - dequeue: remove an element from the front of a queue

Queue operations – cont.

- A currently empty queue that can hold `char` values:




- `enqueue('E') ;`




Queue operations – cont.

- enqueue ('K') ;




- enqueue ('G') ;




Queue operations – cont.



- `dequeue () ; // remove E`



- `dequeue () ; // remove K`



The Queue data structure – Fixed Buffer Implementations



The queue data structure - Design/algorithms


- Create a buffer of length n
- Create some variables : (all initialized to zero)
 - Use a variable (tail or write_index) to hold the index where new data should be written.
 - Use a variable (head or read_index) to point to the array index from which data may be read.
 - A count variable to hold the current number of elements in the array.
- Before enqueueing, make sure that counter < n (i.e. there is room for adding an entry). If not, return a failure.
- To enqueue an entry, write it to the array using the tail variable and then increment the tail using modulo n arithmetic. Also increment the counter

The queue data structure - Design/algorithms

- Before dequeuing, make sure that counter > 0, else return a failure.
- To dequeue an entry, read it out using the head variable, increment the head using modulo n arithmetic and also decrement the counter.


Introduction to the Linked Lists

- Linked list: set of data structures (nodes) that contain references to other data structures
 - References may be addresses or array indices. In our OS class, the kernel uses addresses to point to the next node in the list.
- Nodes may be added to or removed from the linked list during execution



Node Organization

- A node contains:
 - **data**: one or more data fields – may be organized as structure, object, etc.
 - **a pointer**: that can point to another node
- In our case, the data may be the data in the **process control block** (PCB), whereas the pointer is the next PCB (for the next process)



Declaring a Node

- Declare a node:


```
struct ListNode  
{  
    int data;  
    ListNode *next;  
};
```

- Declaraing a struct does not allocate memory

- Instantiating a node allocates memory, e.g.

- `ListNode* ptr = new ListNode;`


Linked List Organization



Linked lists may contain 0 or more nodes:

- Has a list head to point to first node
- Last node points to NULL

An empty list contains 0 nodes,
• The list head points to NULL




Creating a Linked List (an empty one)

- Define a pointer for the head of the list:

```
ListNode *head = nullptr;
```

- Head pointer initialized to `nullptr` to indicate an empty list

- A queue (= FIFO) may be implemented as a linked list.
 - The head of the FIFO is the first entry in the list. Tail is the last entry.




The Null Pointer

- Is used to indicate end-of-list (null = address 0 in memory which normally not a valid address)
- Should always be tested for, before using a pointer:


Linked List Operations

- Basic operations:
 - append a node to the end of the list
 - insert a node within the list
 - traverse the linked list
 - delete a node
 - delete/destroy the list


Linked List operations – appending to tail




Linked List Operations – appending to head



Linked List Operations – inserting a node




Linked List Operations – removing a node



Linked List Queue Operations

- Basic operations:
 - Remove a node from the head
 - Add a node to the tail
 - Queue head and Queue tail are pointers to head node and tail node respectively



Chapter 3: Processes


- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication
- Examples of IPC Systems
- Communication in Client-Server Systems

3.1 The Process Concept

- An operating system executes a variety of programs:
 - Batch system – **jobs**
 - Time-shared systems – **user programs, processes** or **tasks**
- We use the terms ***job, task*** and ***process*** almost interchangeably
 - In conventional full-fledged operating systems, i.e. those running workstations or servers (not embedded systems running FreeRTOS):
process = job = task
- **Process** – a program in execution; process execution generally (but not always) progresses in a sequential fashion.

The Process Concept – cont.


- Area occupied in main memory is divided into multiple sections:
 - The program code, also called **text section**
 - **Data section** containing global variables
 - Initialized sections (aka .data section), followed by
 - Uninitialized sections (aka .bss section)
 - **Stack:** The **call stack** contains temporary data:
 - Saved CPU registers (including return addresses)
 - Function parameters
 - Local variables
 - **Heap** containing memory dynamically allocated during run time. Grows opposite to the stack
 - e.g. using new/delete (in C++ or Java)
 - Malloc/free in C
- A process also occupies/uses CPU registers (not shown on Fig.).




The Call Stack - example

```
#include <iostream>
#include <fstream>
#include <vector>
using namespace std;


int myfunc(int x) {
    int y;
    y = x * x + 2 * x + 5;
    return y;
}
int main() {
    int z;
    z = myfunc(2);
    cout << z << endl;
}
```



Before calling myfunc()



- While myfunc() is executing:
- CPU registers are saved into frame 3
 - Parameter x is pushed into stack frame 3
 - Local variable y is allocated in frame 3




- After myfunc() executed:
- CPU registers are restored from frame 3
 - Parameter x is deallocated
 - Local variable y deallocated
 - Stack frame is no more

The Process Concept (Cont.)

- Program is ***passive*** entity stored on disk (**executable file**), process is ***active***. A program becomes a process when:
 - Its executable file is loaded into main memory and
 - is registered with the scheduler for execution.
- Execution of program is usually started via:
 - GUI mouse clicks
 - Command line entry of its name, etc.
- One program can be instantiated as several processes
 - Consider multiple users executing the same program

Process State

- As a process executes, it changes **state**
 - **new**: The process is being created
 - **ready**: The process is waiting to be assigned to a processor
 - **running**: Instructions are being executed
 - **Waiting**: The process is waiting for some event to occur
 - **terminated**: The process has finished execution




Process Control Block (PCB)

Information associated with each process

(also called **task control block**)

- Process ID
- Process state – running, waiting, etc
- CPU registers – contents of all process-centric registers **including the program counter** (which contains location of next instruction to execute)
- CPU scheduling information - priorities, scheduling queue pointers, etc.
- Memory-management information – memory allocated to the process (base and limit registers, page/segment tables, etc.)
- Accounting information – CPU and real time used, time limits
- Process numbers of parents or children
- Allocated resources – I/O devices allocated to process, list of open files




Threads

- So far, process has a single thread of execution
- Consider having multiple program counters per process
 - Multiple flows of executions == **threads**
 - Running concurrently OR
 - Running in parallel (in case of multiple processor cores)
- Must then have storage for thread details, and multiple program counters in PCB
- More about threads in the next chapter.


Process Representation in Linux

Represented by the C structure `task_struct`

```
pid_t pid; /* process identifier */  
long state; /* state of the process */  
long nice;  
  
unsigned long policy;  
  
struct task_struct *parent; /* this process's parent */  
struct list_head children; /* this process's children */  
struct task_struct *next_task, *prev_task;  
  
struct files_struct *files; /* list of open files */  
struct mm_struct *mm; /* address space of this process */
```



Where is the PCB for process C ?



main memory